

Efficient Software-Based Fault Isolation

Robert Wahbe, Steven Lucco, Thomas E. Anderson,
Susan L. Graham

Problem and Motivation

- Application developers frequently incorporate other code bases into their own applications
 - E.g. it's easier (and usually safer) to use someone else's crypto library than to roll your own
- How can I (an application developer) mitigate negative side effects of third-party code used by my application?

Fault Isolation (FI)

- Isolate components of a system, so that if a fault occurs in one of the isolated components:
 - The side effects of that fault are constrained to the isolated environment
 - It is easy to track down the faulting component and pinpoint the exact cause
- Existing solutions place components in their own address space
 - Depends on hardware to enforce isolation
- Address-space based isolation incurs a significant performance penalty
 - Crossing isolation environments triggers a context switch
 - Context switches are expensive
 - Lots of context switches = lots of overhead

Software Fault Isolation (SFI)

- Load untrusted component into its own memory region
- Ensure that code executed in such a region cannot write or transfer control to memory addresses outside of said memory region
- Advantages
 - Isolation of untrusted components
 - No context switch incurred to communicate between safe-unsafe and unsafe-unsafe, since all sandboxes exist in the same address space

Static or Dynamic Instrumentation

- SFI can be performed both statically and dynamically
- Dynamic
 - Code has already been compiled
 - SFI instrumentation is added via dynamic binary rewriting (“binary patching”)
- Static
 - Compiler-based instrumentation
 - Binaries are generated with SFI enforcement
- Paper chose the static route

Methodology

- Divide an application's virtual address space in segments
- Segments are aligned in such a way that the high bits of an address in that segment are unique to that segment (forms a segment ID)
- Each segment is further split into two regions:
 - Code
 - Data
 - Stack, heap, static data, etc.
- Code is added that checks the target address of writes and control transfers (Segment Matching)?
 - Does the segment ID of the target address match the segment ID of the current segment?

Implementation (Generate Unsafe Instruction Set)

- Use static analysis to reduce the set of instructions that need to be instrumented
 - Possible to statically verify most uses of memory address in write and control transfers
 - Uses of memory derived from offsets relative to the program counter can be checked to be in bounds at compile time
 - Uses of absolute memory addresses
 - E.g. direct calls and jumps
- Indirect writes and control transfers make up the majority of the unsafe instruction set
 - E.g. uses of addresses stored in registers

Implementation (Instrumentation)

`dedicated-reg \leftarrow target address`

Move target address into dedicated register.

`scratch-reg \leftarrow (dedicated-reg \gg shift-reg)`

Right-shift address to get segment identifier.

scratch-reg is not a dedicated register.

shift-reg is a dedicated register.

`compare scratch-reg and segment-reg`

segment-reg is a dedicated register.

`trap if not equal`

Trap if store address is outside of segment.

`store instruction uses dedicated-reg`

Segment Matching

`dedicated-reg \leftarrow target-reg & and-mask-reg`

*Use dedicated register and-mask-reg
to clear segment identifier bits.*

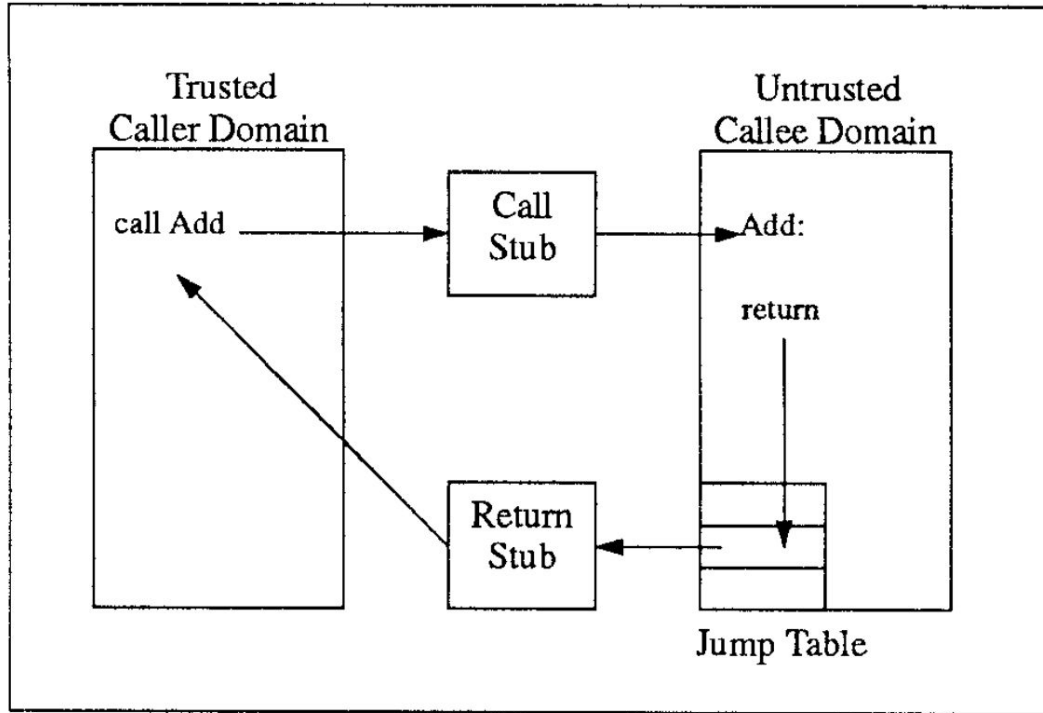
`dedicated-reg \leftarrow dedicated-reg | segment-reg`

*Use dedicated register segment-reg
to set segment identifier bits.*

`store instruction uses dedicated-reg`

Sandboxing

Implementation (Cross-Segment RPC)



- Custom call and return stubs are created for each exported function in the untrusted callee domain
- Destinations in jump table entries are encoded in the instruction
- Jump table is stored in read-only memory
- Call stub copies arguments directly in untrusted callee domain

Implementation (System Resource Usage)

- What if code in an unsafe region needs to access a system resource?
 - E.g. the unsafe code needs to make a syscall
- All accesses to global resources are routed through an “arbitration code” that exists in its own segment
 - Requests for access are passed through cross-domain RPC calls
- If allowed, arbitration code performs the requested action on behalf of the unsafe code, and returns the result (if any)

Evaluation

- Software Encapsulation Overhead
 - Used the sandboxing primitive
 - Benchmarking code was added to an existing code base, but placed inside of an unsafe region
 - ~4.3% on average across all benchmarks
- RPC Overhead
 - Tested with an instrumented version of POSTGRES
 - Overhead measure between ~1.7 and 5.7% across all benchmarks
 - RPC overhead (context switching) estimated to be between ~18.6 and 38.6%

Key Takeaways

- SFI can be used to create isolated regions within the same virtual address space
- SFI is, on average, more performant than FI with hardware-based RPC
 - Author admits that kernel optimizations can be made to reduce the cost of hardware-based RPC, but as long as overhead is above 5%, SFI should perform better

Questions

- There are two SFI implementations described in this paper: binary patching (dynamic) and compiler-enforced (static). Is the dynamic approach possible for x86/64?
 - Disassembling x86/64 is imprecise
 - What about misaligned instructions?

Questions

- If we have to use a compiler-based implementation, is the efficacy of SFI dependant on its adoption rate by writers of third-party libraries?
 - Most developers don't compile their own versions of libraries and carry them around; they typically use shared ones
 - What if a shared library wasn't compiled with SFI enabled?
 - ASLR 2.0?

Questions

- Does the SFI model work if it is dependent on developer adoption?
 - The idea behind SFI is that you, the owner of your code, decide what modules are "unsafe"
 - If library developer adoption is required, then that decision is taken out of the hands of consumers
 - As a library developer, do I have to assume that my code is “unsafe” and incur the performance penalty?